



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# CUDA: Going Parallel with Threads and Kernels

Ben Cumming, CSCS  
July 24, 2016



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Going Parallel : Kernels and Threads

---

## Threads and kernels

- **threads** are streams of execution, run simultaneously on GPU (1000s)
- **kernel** is the task run by each thread
- CUDA provides language support for
  - writing kernels
  - launching many threads to execute a kernel in parallel
- CUDA hides the low-level details of launching threads

## The process for porting to CUDA

1. formulate algorithm in terms of parallel work items
2. write a kernel implementing a work item on one thread
3. launch the kernel with the required number of threads

## Scaled Vector Addition (axpy)

The exercise in the first section used CUBLAS to perform scaled vector addition

$$\mathbf{y} = \mathbf{y} + \alpha \mathbf{x}$$

- $\mathbf{x}$  and  $\mathbf{y}$  are vectors of length  $n$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$$

- $\alpha$  is scalar

$$\alpha \in \mathbb{R}$$

axpy can be performed as  $n$  **independent** operations

$$y_i \leftarrow y_i + a * x_i, \quad i = 0, 1, \dots, n - 1$$

which can be performed independently and in any order

## axpy implemented with for loop

```
void axpy(double *y, double *x, double a, int n) {  
    for(int i=0; i<n; ++i)  
        y[i] = y[i] + a*x[i];  
}
```

## What is a kernel?

- a kernel defines the work item for a single thread
- the work is performed by many threads executing the same kernel **simultaneously**
- Conceptually corresponds to the inner part of a loop for BLAS1 operations like **axpy**

### host : add two vectors

```
void add_cpu(int *a, int *b, int n){  
    for(auto i=0; i<n; ++i)  
        a[i] = a[i] + b[i];  
}
```

### CUDA : add two vectors

```
__global__  
void add_gpu(int *a, int *b, int n){  
    auto i = threadIdx.x;  
    a[i] = a[i] + b[i];  
}
```

- **\_\_global\_\_** keyword indicates a kernel
- **threadIdx** used to find unique id of each thread

## launching a kernel

- host code launches a kernel on the GPU **asynchronously**
- CUDA provides special `<<<_,_>>>` syntax for launching a kernel
  - `add_gpu<<<1, num_threads>>>(args... )` will launch the kernel `add_gpu` with `num_threads` parallel threads.

### host : add two vectors

```
auto n = 1024;
auto a = host_malloc<int>(n);
auto b = host_malloc<int>(n);
add_cpu(a, b, n);
```

### CUDA : add two vectors

```
auto n = 1024;
auto a = device_malloc<int>(n);
auto b = device_malloc<int>(n);
add_gpu<<<1,n>>>(a, b, n);
```

# Exercise: My First Kernel

Open `topics/cuda/practicals/axpy/axpy_kernel.cu`

1. Write a kernel that implements `axpy` for `double`
  - `axpy_kernel(double *y, double *x, double a, int n)`
  - **extra:** can you write a C++ templated version for any type?
2. launch the kernel (look for `TODO`)
3. Compile the test and run
  - it will pass with no errors on success
  - first try with small vectors of size 8
  - try increasing launch size... what happens?
4. **extra:** can you extend the kernel to work for larger arrays?



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Scaling Up : Thread Blocks

---



In the **axpy** exercises we were limited to 1024 threads for a kernel launch

- but we need to scale beyond 1024 threads for the **massive parallelism** we were promised!

## Thread blocks and grids

kernels are executed in groups of threads called **thread blocks**

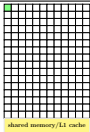
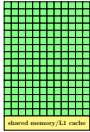
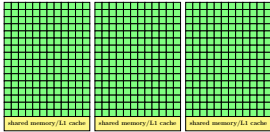
- the launch configuration `axpy<<<grid_dim, block_dim>>>(...)`
  - launch a **grid** of `grid_dim` **blocks**
  - each **block** has `block_dim` **threads**
  - for a total of `grid_dim × block_dim` threads
- previously we launched just one thread block

```
axpy<<<1, n>>>(...)
```

## Why the additional complexity of grids+blocks+threads?

### **Because coordination between threads doesn't scale**

- threads in a block can synchronize and share resources
- this does not scale past a certain number of cores/threads
- on the K20X GPU streaming multiprocessor (SMX) has 192 CUDA cores, and can run 2028 threads
- threads in a block run on the same SMX, with shared resources and thread cooperation
- work is broken into blocks, which are distributed over the 14 SMXs in the K20X GPU

concept	hardware	
thread	 shared memory/L1 cache	<ul style="list-style-type: none"> <li>■ each thread executed on one core</li> </ul>
block	 shared memory/L1 cache	<ul style="list-style-type: none"> <li>■ block executed on 1 SMX</li> <li>■ multiple blocks per SMX if sufficient resources</li> <li>■ threads in a block share SMX resources</li> </ul>
grid	 shared memory/L1 cache   shared memory/L1 cache   shared memory/L1 cache	<ul style="list-style-type: none"> <li>■ kernel is executed in grid of blocks</li> <li>■ blocks distributed over SMXs</li> <li>■ multiple kernels can run at same time</li> </ul>

## Calculating thread indexes

A kernel has to calculate the index of its work item

- in `axpy` we used `threadIdx.x` for the index
- when using multiple blocks, we need more information, which is available in the following **magic variables**:

`gridDim` : total number of blocks in the grid

`blockDim` : number of threads in a thread block

`blockIdx` : index of block `[0, gridDim-1]`

`threadIdx` : index of thread in thread block `[0, blockDim-1]`

## Calculating thread indexes

Consider accessing an array of length 24 with 8 threads per block. The **dimensions** of the kernel launch are:

- `blockDim.x == 8` (8 threads/block)
- `gridDim.x == 3` (3 blocks)

We calculate the index for our thread using the formula

```
auto index = threadIdx.x + blockIdx.x*blockDim.x
```

```
index = threadIdx.x + blockDim.x*blockIdx.x  
      = 5 + 8 * 1  
      = 13
```

threadIdx.x



## Calculating grid dimensions

The number of thread blocks and the number of threads per block are parameters for the kernel launch:

```
kernel<<<blocks, threads_per_block>>>(...)
```

Remember to guard against overflow when the number of work items is not divisible by the thread block size

## vector addition with multiple blocks

```
__global__  
void add_gpu(int *a, int *b, int n){  
    auto i = threadIdx.x + blockIdx.x*blockDim.x;  
    if(i<n) { // guard against access off end of arrays  
        a[i] += b[i];  
    }  
}  
  
// in main()  
auto block_size = 512;  
auto num_blocks = (n + (block_size-1)) / block_size;  
add_gpu<<<num_blocks, block_size>>>(a, b, n);
```

## Calculating grid dimensions

We have to take care when calculating the number of blocks in the grid, i.e. `blocks`:

```
kernel<<<blocks, threads_per_block>>>(...)
```

Most likely, the number of work items `n` is not a multiple of `threads_per_block`.

- in which case some threads in the last thread block will do any work

## Calculating grid dimensions

```
// in main()
auto block_size = 512;
auto num_blocks = (n + (block_size-1)) / block_size;
add_gpu<<<num_blocks, block_size>>>(a, b, n);
```

## The number of threads per block impacts performance

- the optimal number depends on the resources (registers, shared memory, etc) that a kernel requires

## Choosing block size automatically (CUDA 6.5 and later)

```
int block_size, min_grid_dim;  
  
cudaOccupancyMaxPotentialBlockSize(&min_grid_size, &block_size,  
                                   add_gpu, 0, n);  
auto num_blocks = (n + (block_size-1)) / block_size;  
add_gpu<<<num_blocks, block_size>>>(a, b, n);
```

The variable `min_grid_size` is set to the minimum number of blocks required to **saturate** the GPU, i.e. provide the GPU with enough work to utilize all of the SMXs.

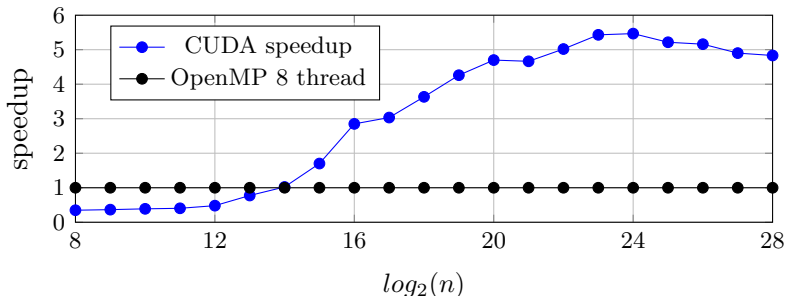


## Exercise: Blocks

Open `cuda/exercises/axpy/axpy.cu` from the last exercise

1. Extend the `axpy` kernel for arbitrarily large input arrays (any `n`)
2. Update the call site to calculate the grid configuration
3. Compile the test and run
  - it will pass with no errors on success
4. Experiment with varying the size of the arrays (scaling)
  - start small and increase
5. finish the `newton.cu` example
  - how do the h2d, d2h and kernel timings compare?
6. **extra:** Compare scaling with the `axpy_omp` benchmark
7. **extra:** Experiment with varying the block size
  - try `cudaOccupancyMaxPotentialBlockSize`.

## Exercise: Results



The GPU is a throughput device:

- the CUDA implementation is faster for  $2^{15} \approx 32,000$
- requires  $2^{20} \approx 1,000,000$  to get “advertised”  $5\times$  speedup

You have to provide enough parallelism to exploit many cores